

This is page one of the test PDF document. It contains sample text for testing the RAG pipeline chunking and retrieval functionality.

This is page two. It discusses vector databases and embedding models used for semantic search in modern applications.